

Bayesian ranking for tennis players in PyMC

PyData Amsterdam 2023

Francesco Bruzzesi

Topics of the day

- What's wrong with the current tennis ranking
- Introduction to Bradley-Terry model
- Implementation in PyMC
- Ranking
- Extensions and other applications

What's wrong with the ranking?



What's wrong with the ranking?

The current system works, but it has a few flaws:

What's wrong with the ranking?

The current system works, but it has a few flaws:

- All opponents are equal

Round	ATP Points
Winner	2 000 points
Finalist	1 200 points
Semi-finalists	720 points
Quarter-finalists	360 points

What's wrong with the ranking?

The current system works, but it has a few flaws:

- All opponents are equal
- It's a number game

name	wins	played	rank	win_rate
Stefanos Tsitsipas	51	71	4	0.718310
Matteo Berrettini	25	34	8	0.735294

What's wrong with the ranking?

The current system works, but it has a few flaws:

- All opponents are equal
- It's a number game
- Surfaces are interchangeable

"I don't want to play here on this (clay) surface" [Daniil Medvedev]

"(...) grass is for golf players" [Casper Ruud]

What's wrong with the ranking?

The current system works, but it has a few flaws:

- All opponents are equal
- It's a number game
- Surfaces are interchangeable
- Last week points count as last year

What's wrong with the ranking?

The current system works, but it has a few flaws:

- All opponents are equal
- It's a number game
- Surfaces are interchangeable
- Last week points count as last year

TL;DR The current ranking system is a *sum* of player performance over the last 52 weeks.

How does Bradley-Terry work?

The Bradley-Terry model is a probabilistic method for *paired comparisons*. It is based on the assumption that the probability of player i beating player j is a function of their abilities ϑ_i and ϑ_j



How does Bradley-Terry work?

The Bradley-Terry model is a probabilistic method for *paired comparisons*. It is based on the assumption that the probability of player i beating player j is a function of their abilities ϑ_i and ϑ_j :

$$P(i \text{ beats } j) = \text{logistic}(\vartheta_i - \vartheta_j)$$



How does Bradley-Terry work?

The Bradley-Terry model is a probabilistic method for *paired comparisons*. It is based on the assumption that the probability of player i beating player j is a function of their abilities ϑ_i and ϑ_j :

$$P(i \text{ beats } j) = \text{logistic}(\vartheta_i - \vartheta_j)$$

We are interested in learning the latent ability ϑ_i for each player from the data (i.e. matches outcome).



Available data

The dataset comes from [Jeff Sackmann github repo](#).

We will focus on main tour from 2021 to 2023 but exclude team, national events, players with less than 10 matches (203 players).

Available data

The dataset comes from [Jeff Sackmann github repo](#).

We will focus on main tour from 2021 to 2023 but exclude team, national events, players with less than 10 matches (203 players).

```
df = load_data(  
    start=2021,  
    end=2023  
)  
# Loaded 6134 matches  
df.tail()
```

Available data

The dataset comes from [Jeff Sackmann github repo](#).

We will focus on main tour from 2021 to 2023 but exclude team, national events, players with less that 10 matches (203 players).

```
df = load_data(  
    start=2021,  
    end=2023  
)  
# Loaded 6134 matches  
df.tail()
```

Tourney	Date	Surface	Winner Name	Loser Name	Winner Rank	Loser Rank
Hamburg	2023-07-24	Clay	A. Zverev	A. Fils	19	71
Umag	2023-07-24	Clay	A. Popyrin	M. Arnaldi	90	76
Atlanta	2023-07-24	Hard	T. Fritz	A. Vukic	9	82
Hamburg	2023-07-24	Clay	A. Zverev	L. Djere	19	57
Umag	2023-07-24	Clay	A. Popyrin	S. Wawrinka	90	72

Cross Validation

Given a week to forecast, we train with observations one year prior.

Time series cross validation



Baseline

A naive model assumes that the player with the better ranking will win the match. This will establish a baseline for the “predictive power” of the current ranking system.

Baseline

A naive model assumes that the player with the better ranking will win the match. This will establish a baseline for the “predictive power” of the current ranking system.

```
(df
    .assign(best_rank_win = lambda t: t["winner_rank"]<t["loser_rank"]))
```

Baseline

A naive model assumes that the player with the better ranking will win the match. This will establish a baseline for the “predictive power” of the current ranking system.

```
(df
    .assign(best_rank_win = lambda t: t["winner_rank"]<t["loser_rank"])
    .loc[lambda t: t["year"].ge(2022), "best_rank_win"]
)
```

Baseline

A naive model assumes that the player with the better ranking will win the match. This will establish a baseline for the “predictive power” of the current ranking system.

```
(df
  .assign(best_rank_win = lambda t: t["winner_rank"]<t["loser_rank"])
  .loc[lambda t: t["year"].ge(2022), "best_rank_win"]
  .mean()
)
```

```
# 0.618
```

“Ability Based” model

“Ability Based” model

```
import pymc as pm
```

```
with pm.Model() as base_model:
```

“Ability Based” model

```
import pymc as pm
```

```
with pm.Model() as base_model:
```

```
    X, y = pm.MutableData("X", X_train), pm.MutableData("y", y_train)  
    player1, player2 = X[:, 0], X[:, 1]
```

sample values

```
X = array([[  
    63, 47],  
    90, 46],  
    ...,  
    [89, 61]  
]) # shape=(n_matches, 2)  
y = array([1, 1, ..., 1]) # shape=(n_matches, )
```

“Ability Based” model

```
import pymc as pm
```

```
with pm.Model() as base_model:
```

```
    X, y = pm.MutableData("X", X_train), pm.MutableData("y", y_train)  
    player1, player2 = X[:, 0], X[:, 1]
```

```
    a_m = pm.Normal("ability_m", mu=0.0, sigma=1., shape=(n_players_,))  
    a_sd = pm.HalfCauchy("ability_sd", beta=1.0)
```

sample values

```
a_m = array([0.5, -0.4, ..., 0.1]) # shape=(n_players, )  
a_s = array([2.0]) # shape=(1, )
```


“Ability Based” model

```
import pymc as pm

with pm.Model() as base_model:

    X, y = pm.MutableData("X", X_train), pm.MutableData("y", y_train)
    player1, player2 = X[:, 0], X[:, 1]

    a_m = pm.Normal("ability_m", mu=0.0, sigma=1., shape=(n_players,))
    a_sd = pm.HalfCauchy("ability_sd", beta=1.0)

    player_ability = pm.Deterministic("player_ability", a_m*a_sd)
    delta_ability = player_ability[player1] - player_ability[player2]
```

sample values

```
player_ability = array([3.3, -0.7, ..., 1.1]) # shape=(n_players, )
delta_ability = array([2.1, -1.2, ..., 0.2]) # shape=(n_matches, )
```

“Ability Based” model

```
import pymc as pm

with pm.Model() as base_model:

    X, y = pm.MutableData("X", X_train), pm.MutableData("y", y_train)
    player1, player2 = X[:, 0], X[:, 1]

    a_m = pm.Normal("ability_m", mu=0.0, sigma=1., shape=(n_players_,))
    a_sd = pm.HalfCauchy("ability_sd", beta=1.0)

    player_ability = pm.Deterministic("player_ability", a_m*a_sd)
    delta_ability = player_ability[player1] - player_ability[player2]

    prob = pm.Deterministic("prob", pm.invlogit(delta_ability))
    _ = pm.Bernoulli("result", p=prob, observed=y)
```

sample values

```
prob = array([0.91, 0.87, ..., 0.38])
result = array([1, 1, ..., 0]) # shape=(n_matches, )
```

Fit the model

In order to fit the model, we use the *inference magic button*.


```
with base_model:
```

Fit the model

In order to fit the model, we use the *inference magic button*.

```
with base_model:
    base_trace = pm.sample(
        draws=1000,
        tune=1000,
        chains=4,
        nuts_sampler="numpyro",
        nuts_sampler_kwargs={"chain_method": "parallel"},
        ...
    )
```

Fit the model



Francesco Bruzzesi

Posterior trace

Trace contains all sample stats and posterior information

Posterior trace

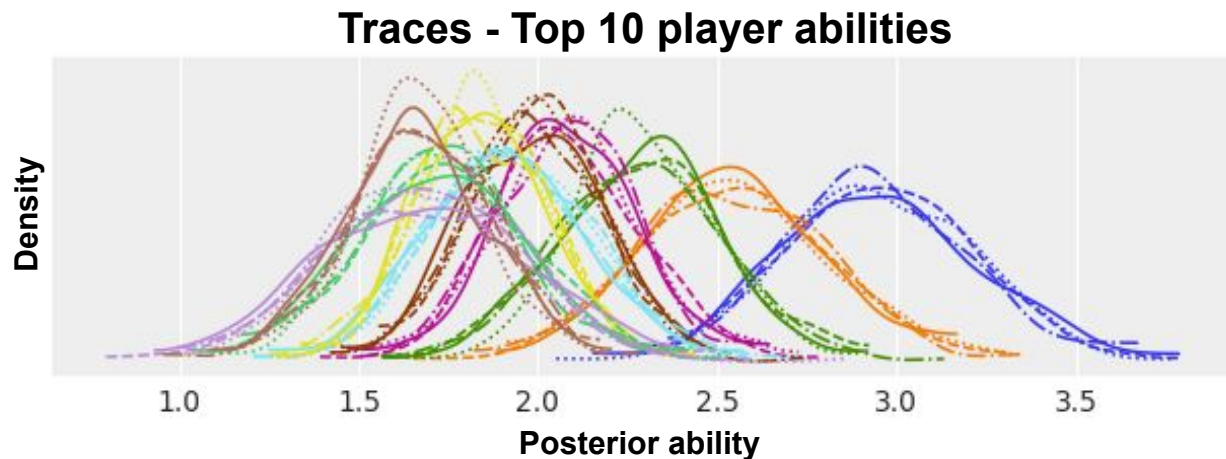
Trace contains all sample stats and posterior information

base_trace
✓ 0.3s
arviz.InferenceData
▸ posterior
▸ sample_stats
▸ observed_data
▸ constant_data

Posterior trace

Trace contains all sample stats and posterior information

base_trace
✓ 0.3s
arviz.InferenceData
► posterior
► sample_stats
► observed_data
► constant_data



Predict on new data

Finally we can predict on out of sample data

Predict on new data

Finally we can predict on out of sample data

```
with base_model:  
    test_size = X_test.shape[0]  
    pm.set_data({  
        "X": X_test,  
        "y": np.empty(test_size, dtype=int)  
    })
```

Predict on new data

Finally we can predict on out of sample data

```
with base_model:
    test_size = X_test.shape[0]
    pm.set_data({
        "X": X_test,
        "y": np.empty(test_size, dtype=int)
    })

    base_posterior = pm.sample_posterior_predictive(
        trace=base_trace, var_names=[...]
    )
```

Predict on new data

Finally we can predict on out of sample data

```
with base_model:
    test_size = X_test.shape[0]
    pm.set_data({
        "X": X_test,
        "y": np.empty(test_size, dtype=int)
    })

    base_posterior = pm.sample_posterior_predictive(
        trace=base_trace, var_names=[...]
    )
```

base_posterior

✓ 0.0s

arviz.InferenceData

► posterior_predictive

► constant_data

Predict on new data

Finally we can predict on out of sample data

```
with base_model:
    test_size = X_test.shape[0]
    pm.set_data({
        "X": X_test,
        "y": np.empty(test_size, dtype=int)
    })

    base_posterior = pm.sample_posterior_predictive(
        trace=base_trace, var_names=[...]
    )
```

base_posterior

✓ 0.0s

arviz.InferenceData

► posterior_predictive

► constant_data

Model performance: 62.3% accuracy (0.5% above *baseline*)

“Ability, by Surface, with Decay”

“Ability, by Surface, with Decay”

```
with pm.Model() as model:
```

```
    X, y = pm.MutableData("X", X), pm.MutableData("y", y)
```

```
    player1, player2 = X[:, 0].astype(int), X[:, 1].astype(int)
```

```
    surface, sample_weights = X[:, 2].astype(int), X[:, 3]
```

“Ability, by Surface, with Decay”

```
with pm.Model() as model:
```

```
    X, y = pm.MutableData("X", X), pm.MutableData("y", y)
```

```
    player1, player2 = X[:, 0].astype(int), X[:, 1].astype(int)
```

```
    surface, sample_weights = X[:, 2].astype(int), X[:, 3]
```

```
    ability_m = pm.Normal("ability_m", 0.0, 1.0, shape=(n_players_,))
```

```
    ability_sd = pm.HalfCauchy("ability_sd", beta=1.0)
```

```
    base_ability = pm.Deterministic("base_ability", ability_m * ability_sd)
```

```
    base_delta = player_ability[player1] - player_ability[player2]
```


“Ability, by Surface, with Decay”

```
with pm.Model() as model:

    X, y = pm.MutableData("X", X), pm.MutableData("y", y)

    player1, player2 = X[:, 0].astype(int), X[:, 1].astype(int)
    surface, sample_weights = X[:, 2].astype(int), X[:, 3]

    ability_m = pm.Normal("ability_m", 0.0, 1.0, shape=(n_players_,))
    ability_sd = pm.HalfCauchy("ability_sd", beta=1.0)
    base_ability = pm.Deterministic("base_ability", ability_m * ability_sd)
    base_delta = player_ability[player1] - player_ability[player2]

    surface_m = pm.Normal("surface_m", 0.0, 1.0, shape=(n_players_, n_surfaces_))
    surface_sd = pm.HalfCauchy("surface_sd", beta=1.0)
    surface_factor = pm.Deterministic("surface_factor", surface_m * surface_sd)
    surface_delta = player_surface[player1, surface] - player_surface[player2, surface]
```

“Ability, by Surface, with Decay”

```
with pm.Model() as model:

    X, y = pm.MutableData("X", X), pm.MutableData("y", y)

    player1, player2 = X[:, 0].astype(int), X[:, 1].astype(int)
    surface, sample_weights = X[:, 2].astype(int), X[:, 3]

    ability_m = pm.Normal("ability_m", 0.0, 1.0, shape=(n_players_,))
    ability_sd = pm.HalfCauchy("ability_sd", beta=1.0)
    base_ability = pm.Deterministic("base_ability", ability_m * ability_sd)
    base_delta = player_ability[player1] - player_ability[player2]

    surface_m = pm.Normal("surface_m", 0.0, 1.0, shape=(n_players_, n_surfaces_))
    surface_sd = pm.HalfCauchy("surface_sd", beta=1.0)
    surface_factor = pm.Deterministic("surface_factor", surface_m * surface_sd)
    surface_delta = player_surface[player1, surface] - player_surface[player2, surface]

    prob = pm.Deterministic("prob", pm.invlogit(delta_ability + delta_surface))
    logp = sample_weights * pm.logp(pm.Bernoulli.dist(p=prob), y)
    _ = pm.Potential("error", logp)
```

“Ability, by Surface, with Decay”

As before, we proceed to:

- Fit the model
- Check there are no issues in convergence
- Run the full backtest/cross validation

Model performance: 63.2% accuracy (1.4% above *baseline*)

“Ability, by Surface, with Decay”



Ranking

Ranking(s) resulting from the model is not too different when compared to the actual one, with a few *big* adjustments

Ranking

Ranking(s) resulting from the model is not too different when compared to the actual one, with a few *big* adjustments

Ranking at 2023-07-24

player_name	rank	core_rank	hard_rank	clay_rank	grass_rank	core_ability	hard_ability	clay_ability	grass_ability
Carlos Alcaraz	1	2	2	2	1	2.730000	2.700000	2.730000	2.960000
Novak Djokovic	2	1	1	1	2	2.890000	3.030000	2.850000	2.930000
Daniil Medvedev	3	3	3	3	3	2.190000	2.420000	2.170000	2.130000
Casper Ruud	4	16	23	8	20	0.730000	0.690000	1.240000	0.690000
Stefanos Tsitsipas	5	7	9	7	7	1.290000	1.380000	1.450000	1.180000
Holger Rune	6	6	7	6	6	1.480000	1.480000	1.590000	1.530000
Andrey Rublev	7	5	6	5	5	1.510000	1.500000	1.610000	1.570000
Jannik Sinner	8	4	4	4	4	1.840000	1.910000	1.870000	1.880000
Taylor Fritz	9	13	10	14	15	0.890000	1.280000	0.920000	0.790000
Frances Tiafoe	10	8	12	11	8	1.100000	1.190000	1.090000	1.180000

Ranking

Ranking(s) resulting from the model is not too different when compared to the actual one, with a few *big* adjustments

Ranking at 2023-07-24

player_name	rank	core_rank	hard_rank	clay_rank	grass_rank	core_ability	hard_ability	clay_ability	grass_ability
Carlos Alcaraz	1	2	2	2	1	2.730000	2.700000	2.730000	2.960000
Novak Djokovic	2	1	1	1	2	2.890000	3.030000	2.850000	2.930000
Daniil Medvedev	3	3	3	3	3	2.190000	2.420000	2.170000	2.130000
Casper Ruud	4	16	23	8	20	0.730000	0.690000	1.240000	0.690000
Stefanos Tsitsipas	5	7	9	7	7	1.290000	1.380000	1.450000	1.180000
Holger Rune	6	6	7	6	6	1.480000	1.480000	1.590000	1.530000
Andrey Rublev	7	5	6	5	5	1.510000	1.500000	1.610000	1.570000
Jannik Sinner	8	4	4	4	4	1.840000	1.910000	1.870000	1.880000
Taylor Fritz	9	13	10	14	15	0.890000	1.280000	0.920000	0.790000
Frances Tiafoe	10	8	12	11	8	1.100000	1.190000	1.090000	1.180000



Ranking

Ranking(s) resulting from the model is not too different when compared to the actual one, with a few *big* adjustments

Ranking at 2023-07-24

player_name	rank	core_rank	hard_rank	clay_rank	grass_rank	core_ability	hard_ability	clay_ability	grass_ability
Carlos Alcaraz	1	2	2	2	1	2.730000	2.700000	2.730000	2.960000
Novak Djokovic	2	1	1	1	2	2.890000	3.030000	2.850000	2.930000
Daniil Medvedev	3	3	3	3	3	2.190000	2.420000	2.170000	2.130000
Casper Ruud	4	16	23	8	20	0.730000	0.690000	1.240000	0.690000
Stefanos Tsitsipas	5	7	9	7	7	1.290000	1.380000	1.450000	1.180000
Holger Rune	6	6	7	6	6	1.480000	1.480000	1.590000	1.530000
Andrey Rublev	7	5	6	5	5	1.510000	1.500000	1.610000	1.570000
Jannik Sinner	8	4	4	4	4	1.840000	1.910000	1.870000	1.880000
Taylor Fritz	9	13	10	14	15	0.890000	1.280000	0.920000	0.790000
Frances Tiafoe	10	8	12	11	8	1.100000	1.190000	1.090000	1.180000

Comparison over time




Why Bayesian?




Why Bayesian?



Why Bayesian?



Why Bayesian?



Extensions & other applications

It's possible to:

- Include other factors (e.g. fatigue)
- Consider priors more robust to outliers (e.g. Student T)
- Use a different model architecture (e.g. a hierarchical model, gaussian processes)
- Extend the period of analysis
- Compare with other models (e.g. ELO Rating)

Where to find me?

